

CS 553 Cloud Computing

Programming Assignment 1

SOURCE CODE

Dinesh Chowday Jasti (A20335710)

Jayanth Vangari (A20337867)

Sai Sravan Rachiraju (A20332891)

NormalJavaWCount.java:

```
import java.io.*;

import java.util.*;

import java.util.concurrent.Semaphore;


public class NormalJavaWCount {

    private static int ThreadCount = 1;

    private static int FileCount = 16;

    private static List<String> fileList = new ArrayList<String>();

    private static HashMap<String, Integer> hmap = new HashMap<String, Integer>();

    private static String filename = null;


    @SuppressWarnings({ "unchecked" })

    public static void main(String[] args) throws InterruptedException,

        IOException {

        BufferedReader input = new BufferedReader(new InputStreamReader(

            System.in));

        System.out.println("Enter the File Name: ");

        filename = input.readLine().trim();


        //filesplitter(filename);

        fileListLoader();


        System.out.println("Enter number of Threads: ");

        try {

            ThreadCount = Integer.parseInt(input.readLine().trim());
```

```

    } catch (Exception e) {
    }

    Thread[] threadnumber = new Thread[ThreadCount];

    final Semaphore semFiles = new Semaphore(ThreadCount, false);

    long start = System.currentTimeMillis();

    for (int i = 0; i < ThreadCount; i++) {
        Thread.sleep(i * 10);

        threadnumber[i] = new Thread(new Runnable() {

            @Override

            public void run() {

                try {

                    semFiles.acquire();

                    while (!fileList.isEmpty()) {

                        String filename = fileList.remove(0);

                        Wordcount(filename);

                    }

                    semFiles.release();

                } catch (Exception ex) {

                }

            }

        });

        threadnumber[i].start();
    }

    for (int i = 0; i < ThreadCount; i++) {

        threadnumber[i].join();

    }

```

```

        long end = System.currentTimeMillis();

        float diff = end - start;

        System.out.println("\nTotal time taken: " + diff + " milliseconds");

        System.out.println("Execution Time(s): " + diff/1000);

        System.out.println("Execution Time(min): " + diff / (1000 *60 ));

        System.out.println("Execution Time(hr): " + diff / (1000 * 60 * 60) + "\n");


        FileWriter fwstream = new FileWriter("wordcount-java.txt");

        System.out.println("The total number of unique words are: " + hmap.size());

        BufferedWriter out = new BufferedWriter(fwstream);


        for(Map.Entry<String,Integer> entry : hmap.entrySet()) {

            String key = entry.getKey();

            Integer value = entry.getValue();

            out.write("%20s".format(key) + ": ");

            out.write(" %10d".format(Integer.toString(value)));

            out.newLine();

        }


        out.close();

    }


    public static void fileListLoader() {

        String name = "file";

        for (int i = 1; i <= FileCount; i++) {

            if (i >= 10)

```

```

        name = "file";

        fileList.add(name + i + ".txt");

    }

}

public static void createFilesChunk(int start, long end, String fname) {

    File file = new File(filename);

    try {

        FileReader reader = new FileReader(file);

        LineNumberReader lreader = new LineNumberReader(reader);

        String lines = "";

        FileWriter fWriter = new FileWriter(fname);

        BufferedWriter bWriter = new BufferedWriter(fWriter);

        while (lreader.getLineNumber() != start) {

            lines = lreader.readLine();

        }

        lreader.setLineNumber(start);

        while (lreader.getLineNumber() != end) {

            lines = lreader.readLine();

            bWriter.write(lines);

            bWriter.newLine();

        }

        bWriter.close();

        lreader.close();

        System.out.println(fname + ": Starting Line: " + start

            + "          Ending Line: " + end);
    }
}

```

```

        } catch (Exception ex) {
        }
    }
}

```

```

private static Semaphore sema = new Semaphore(ThreadCount, false);

```

```

public static void Wordcount(String filename) {
    String str;
    try {
        System.out.println(Thread.currentThread().getName()
            + " ---> assigned to ---> " + filename);
        BufferedReader br = new BufferedReader(new FileReader(filename));
        while ((str = br.readLine()) != null) {
            str = str.replaceAll("[^a-zA-Z0-9\\s]", "");
            StringTokenizer stringTok = new StringTokenizer(str);
            int count = stringTok.countTokens();
            for (int l = 0; l < count; l++) {
                String word = stringTok.nextToken();
                if(word.equals(""))
                {
                    continue;
                }
                else
                {
                    sema.acquire();
                    hmap.put(word, hmap.get(word) == null ? 1

```

```

                                : ((Integer) hmap.get(word) + 1));
                                sema.release();
                                }
                                }
                                }
        } catch (Exception ex) {
            ex.printStackTrace();
        }
    }
}

```

```

private static LineNumberReader lreader;

```

```

public static int TotalLines(String fileName) {
    int total_lines = 0;
    try {
        File f = new File(fileName);
        FileReader reader = new FileReader(f);
        lreader = new LineNumberReader(reader);
        while ((lreader.readLine()) != null) {
        }
        total_lines = lreader.getLineNumber();
    } catch (Exception ex) {
    }
    return total_lines;
}

```

```
public static void filesplitter(String fileName) {  
  
    int start = 0;  
  
    int total_lines = 123015884;  
  
    int diff = total_lines / FileCount;  
  
    int end = diff;  
  
    String a;  
  
    for (int i = 1; i <= FileCount; i++) {  
  
        a = (i < 10 ? "file" : "file") + i + ".txt";  
  
        createFilesChunk(start, end, a);  
  
        start = end;  
  
        end = end + diff;  
  
        if ((i == (FileCount - 1)) && (end != total_lines))  
            end = total_lines;  
  
    }  
  
}  
  
}
```


SortedJavaWCount.java

```
import java.io.*;

import java.util.*;

import java.util.concurrent.Semaphore;


public class SortedJavaWCount {

    private static int ThreadCount = 1;

    private static int FileCount = 16;

    private static List<String> fileList = new ArrayList<String>();

    private static TreeMap<String, Integer> hmap = new TreeMap<String, Integer>();

    private static String filename = null;


    @SuppressWarnings({ "unchecked" })

    public static void main(String[] args) throws InterruptedException,

        IOException {

        BufferedReader input = new BufferedReader(new InputStreamReader(

            System.in));

        System.out.println("Enter the File Name: ");

        filename = input.readLine().trim();


        //filesplitter(filename);

        fileListLoader();


        System.out.println("Enter number of Threads: ");
```

```

try {
    ThreadCount = Integer.parseInt(input.readLine().trim());
} catch (Exception e) {
}

Thread[] threadnumber = new Thread[ThreadCount];

final Semaphore semFiles = new Semaphore(ThreadCount, false);

long start = System.currentTimeMillis();

for (int i = 0; i < ThreadCount; i++) {
    Thread.sleep(i * 10);

    threadnumber[i] = new Thread(new Runnable() {

        @Override
        public void run() {

            try {

                semFiles.acquire();

                while (!fileList.isEmpty()) {

                    String filename = fileList.remove(0);

                    Wordcount(filename);

                }

                semFiles.release();

            } catch (Exception ex) {

            }

        }

    });

    threadnumber[i].start();
}

for (int i = 0; i < ThreadCount; i++) {

```

```

        threadnumber[i].join();
    }

    FileWriter fwstream = new FileWriter("sort1mb-java.txt");

    System.out.println("The total number of unique words are: " + hmap.size());

    BufferedWriter out = new BufferedWriter(fwstream);

        out.write("Sorted word count without frequency based on the ASCII");

    for(Map.Entry<String,Integer> entry : hmap.entrySet()) {

        String key = entry.getKey();

        Integer value = entry.getValue();

        out.write("%20s".format(key));

        out.newLine();

    }

    out.close();

    long end = System.currentTimeMillis();

    float diff = end - start;

    float splitdiff = 1348032;

    float total_time=splitdiff+diff;

    System.out.println("\nTotal time taken: " + total_time + " milliseconds");

    System.out.println("Execution Time(s): " + total_time/1000);

    System.out.println("Execution Time(min): " + total_time / (1000 *60 ));

    System.out.println("Execution Time(hr): " + total_time / (1000 * 60 * 60) + "\n");

}

```

```

public static void fileListLoader() {

    String name = "file";

    for (int i = 1; i <= FileCount; i++) {

        if (i >= 10)

            name = "file";

        fileList.add(name + i + ".txt");

    }

}

public static void createFilesChunk(int start, long end, String fname) {

    File file = new File(filename);

    try {

        FileReader reader = new FileReader(file);

        LineNumberReader lreader = new LineNumberReader(reader);

        String lines = "";

        FileWriter fWriter = new FileWriter(fname);

        BufferedWriter bWriter = new BufferedWriter(fWriter);

        while (lreader.getLineNumber() != start) {

            lines = lreader.readLine();

        }

        lreader.setLineNumber(start);

        while (lreader.getLineNumber() != end) {

            lines = lreader.readLine();

            bWriter.write(lines);

            bWriter.newLine();

        }

    }

}

```

```

        bWriter.close();

        lreader.close();

        System.out.println(fname + ": Starting Line: " + start
                                + "          Ending Line: " + end);
    } catch (Exception ex) {
    }
}

private static Semaphore sema = new Semaphore(ThreadCount, false);

public static void Wordcount(String filename) {
    String str;
    try {
        System.out.println(Thread.currentThread().getName()
                                + " ---> assigned to ---> " + filename);
        BufferedReader br = new BufferedReader(new FileReader(filename));
        while ((str = br.readLine()) != null) {
            str = str.replaceAll("[^a-zA-Z0-9\\s]", "");
            StringTokenizer stringTok = new StringTokenizer(str);
            int count = stringTok.countTokens();
            for (int l = 0; l < count; l++) {
                String word = stringTok.nextToken();
                if(word.equals(""))
                {
                    continue;
                }
            }
        }
    }
}

```

```

        else
        {
            sema.acquire();

            hmap.put(word, hmap.get(word) == null ? 1
                        : ((Integer) hmap.get(word) + 1));

            sema.release();
        }
    }
} catch (Exception ex) {
    ex.printStackTrace();
}
}

```

```

private static LineNumberReader lreader;

```

```

public static int TotalLines(String fileName) {
    int total_lines = 0;
    try {
        File f = new File(fileName);
        FileReader reader = new FileReader(f);
        lreader = new LineNumberReader(reader);
        while ((lreader.readLine()) != null) {
        }
        total_lines = lreader.getLineNumber();
    } catch (Exception ex) {
    }
}

```

```
    }

    return total_lines;
}

public static void filesplitter(String fileName) {
    int start = 0;
    int total_lines = TotalLines(fileName);
    int diff = total_lines / FileCount;
    int end = diff;
    String a;
    for (int i = 1; i <= FileCount; i++) {
        a = (i < 10 ? "file" : "file") + i + ".txt";
        createFilesChunk(start, end, a);
        start = end;
        end = end + diff;
        if ((i == (FileCount - 1)) && (end != total_lines))
            end = total_lines;
    }
}
}
```

Hadoop Word Count:

Wc.jar

```
import java.io.IOException;

import java.util.*;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.*;

import org.apache.hadoop.mapred.*;

import org.apache.hadoop.mapreduce.Job;


public class WordCount {

    public static class Map extends MapReduceBase implements Mapper<LongWritable, Text,
Text, IntWritable>
    {

        private final static IntWritable one = new IntWritable(1);

        private Text word = new Text();


        public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable>
output, Reporter reporter) throws IOException
        {

            String line = value.toString();

            String temp;

            StringTokenizer tokenizer = new StringTokenizer(line);

            while (tokenizer.hasMoreTokens())
            {
```



```
        temp = tokenizer.nextToken();

        temp = temp.replaceAll("[^a-zA-Z0-9\\s]", "");

        word.set(temp);

        output.collect(word, one);
    }
}
}
```

```
public static class Reduce extends MapReduceBase implements Reducer<Text, IntWritable,
Text, IntWritable>
```

```
{
```

```
    public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable>
output, Reporter reporter) throws IOException
```

```
    {
```

```
        int sum = 0;
```

```
        while (values.hasNext())
```

```
        {
```

```
            sum += values.next().get();
```

```
        }
```

```
        output.collect(key, new IntWritable(sum));
```

```
    }
```

```
}
```

```
public static class Map_sort extends MapReduceBase implements Mapper<Object, Text,
IntWritable, Text>
```

```
{
```

```
    public void map(Object key, Text value, OutputCollector<IntWritable, Text> collector,
Reporter arg3) throws IOException
```

```
    {
```

```
        String line = value.toString();
```

```
        StringTokenizer stringTokenizer = new StringTokenizer(line);
```

```
        {
```

```
            int number=0;
```

```
            String word=null;
```

```
            if (stringTokenizer.hasMoreTokens())
```

```
                {
```

```
                    String str0 = stringTokenizer.nextToken();
```

```
                    word = str0.trim();
```

```
                }
```

```
            if (stringTokenizer.hasMoreElements())
```

```
                {
```

```
                    String str1 = stringTokenizer.nextToken();
```

```
                    number = Integer.parseInt(str1.trim());
```

```
                }
```

```
            collector.collect(new IntWritable(number), new Text(word));
```

```
        }
```

```
}
```

```
}
```

```
public static class Reduce_sort extends MapReduceBase implements Reducer<IntWritable,  
Text, IntWritable, Text>
```

```
{
```

```
    public void reduce(IntWritable key, Iterator<Text> values, OutputCollector<IntWritable,  
Text> arg2, Reporter arg3) throws IOException
```

```
    {
```

```
        while ((values.hasNext()))
```

```
        {
```

```
            arg2.collect(key, values.next());
```

```
        }
```

```
    }
```

```
}
```

```
public static void main(String[] args) throws Exception {
```

```
    JobConf conf = new JobConf(WordCount.class);
```

```
    conf.setJobName("wordCount");
```

```
    conf.setOutputKeyClass(Text.class);
```

```
    conf.setOutputValueClass(IntWritable.class);
```

```
conf.setMapperClass(Map.class);
conf.setCombinerClass(Reduce.class);
conf.setReducerClass(Reduce.class);

conf.setInputFormat(TextInputFormat.class);
conf.setOutputFormat(TextOutputFormat.class);

FileInputFormat.setInputPaths(conf, new Path(args[0]));
FileOutputFormat.setOutputPath(conf, new Path("/tmp/temp"));

Job job1 = new Job(conf);
job1.submit();
//JobClient.runJob(conf);

//-----
JobConf conf2 = new JobConf(WordCount.class);
conf2.setJobName("WordCount1 ");

conf2.setOutputKeyClass(IntWritable.class);
conf2.setOutputValueClass(Text.class);

conf2.setMapperClass(Map_sort.class);
conf2.setCombinerClass(Reduce_sort.class);
conf2.setReducerClass(Reduce_sort.class);
```

```
conf2.setInputFormat(TextInputFormat.class);  
conf2.setOutputFormat(TextOutputFormat.class);
```

```
FileInputFormat.setInputPaths(conf2, new Path("/tmp/temp/part-00000"));  
FileOutputFormat.setOutputPath(conf2, new Path(args[1]));
```

```
Job job2 = new Job(conf2);  
if (job1.waitForCompletion(true))  
{  
    job2.submit();  
    job2.waitForCompletion(true);  
}  
  
}  
  
}
```

Ws.jar: for sorted hadoop:

```
import java.io.IOException;

import java.util.*;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.*;

import org.apache.hadoop.mapred.*;

import org.apache.hadoop.mapreduce.Job;


public class WordSort {

    public static class Map extends MapReduceBase implements Mapper<LongWritable, Text,
Text, Text>

    {

        private Text word = new Text();


        public void map(LongWritable key, Text value, OutputCollector<Text, Text> output,
Reporter reporter) throws IOException

        {

            String line = value.toString();

            String temp;

            StringTokenizer tokenizer = new StringTokenizer(line);

            while (tokenizer.hasMoreTokens())

            {

                temp = tokenizer.nextToken();

                temp = temp.replaceAll("[^a-zA-Z0-9\\s]", "");

                word.set(temp);
```

```
        output.collect(word, new Text(""));
    }
}
}
```

```
public static class Reduce extends MapReduceBase implements Reducer<Text, Text, Text,
Text>
```

```
{
```

```
    public void reduce(Text key, Iterator<Text> values, OutputCollector<Text, Text> output,
Reporter reporter) throws IOException
```

```
    {
        output.collect(key, new Text(""));
    }
}
```

```
public static void main(String[] args) throws Exception {
```

```
    JobConf conf = new JobConf(WordSort.class);
```

```
    conf.setJobName("WordSort");
```

```
    conf.setOutputKeyClass(Text.class);
```

```
    conf.setOutputValueClass(Text.class);
```

```
    conf.setMapperClass(Map.class);
```

```
    conf.setCombinerClass(Reduce.class);
```

```
    conf.setReducerClass(Reduce.class);
```

```
conf.setInputFormat(TextInputFormat.class);  
conf.setOutputFormat(TextOutputFormat.class);  
  
FileInputFormat.setInputPaths(conf, new Path(args[0]));  
FileOutputFormat.setOutputPath(conf, new Path(args[1]));  
  
Job job1 = new Job(conf);  
job1.submit();  
job1.waitForCompletion(true);  
}  
}
```


Configuration files hadoop:

The following files will have to be modified to complete the Hadoop setup:

- ~/.bashrc
- / conf /hadoop-env.sh
- / conf /core-site.xml
- / conf /yarn-site.xml
- / conf /hdfs-site.xml
- / conf /mapred-site.xml.template

1. ~/.bashrc:

```
#HADOOP VARIABLES START
export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64
export HADOOP_INSTALL=/usr/ubuntu/installs/hadoop
export PATH=$PATH:$HADOOP_INSTALL/bin
export PATH=$PATH:$HADOOP_INSTALL/sbin
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_HDFS_HOME=$HADOOP_INSTALL
export YARN_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native
export HADOOP_OPTS="-Djava.library.path=$HADOOP_INSTALL/lib"
eval ($ssh-agent);
ssh-add/home/ubuntu/dj_key.pem
#HADOOP VARIABLES END
```

Reboot the connection now for the changes to be reflected.

2. / conf /hadoop-env.sh:

This is to set the JAVA Path by modifying **hadoop-env.sh** file

```
export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64
```

3. / conf /core-site.xml:

The **/usr/local/hadoop/etc/hadoop/core-site.xml** file contains configuration properties that Hadoop uses when starting up. This file can be used to override the default settings that Hadoop starts with.

```
<configuration>
<property>
<name>fs.default.name</name>
<value>hdfs://<172.31.22.90>:9000</value>
</property>
<property>
<name>hadoop.tmp.dir</name>
<value>/home/ubuntu/installs/hadoop/hadoop_tmp/tmp</value>
</property>
</configuration>
```

4. / conf /yarn-site.xml:

```
<configuration>
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
<value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
<property>
<name>yarn.resourcemanager.resource-tracker.address</name>
<value><172.31.22.90>:8025</value>
</property>
<property>
<name>yarn.resourcemanager.scheduler.address</name>
<value><172.31.22.90>:8030</value>
</property>
<property>
<name>yarn.resourcemanager.address</name>
<value><172.31.22.90>:8040</value>
</property>
</configuration>
```

5. / conf /hdfs-site.xml:

```
<configuration>
<property>
<name>dfs.replication</name>
<value>16</value>
</property>
<property>
<name>dfs.permissions</name>
<value>>false</value>
</property>
</configuration>
```

6. /conf /mapred-site.xml:

By default, the /usr/local/hadoop/etc/hadoop/ contains the /usr/local/hadoop/etc/hadoop/mapred-site.xml.template file which has to be renamed/copied with the name mapred-site.xml:

“cp /usr/local/hadoop/etc/hadoop/mapred-site.xml.template/usr/local/hadoop/etc/hadoop/mapred-site.xml”

The **mapred-site.xml** file is used to specify which framework is being used for MapReduce. We need to enter the following content in between the <configuration></configuration> tag:

```
<configuration>
<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>
</configuration>
```

Swift:

regexexpressionmapper.swift:

```
type filee;

type count;

app (count t) countwords (filee f) {

wc "-w" @filename(f) stdout=@filename(t);}

messagefile inputfile <"regexexpressionmapper.words.txt">;

countfile c <regexexpressionmapper;
source=@inputfile,match="(.)txt",transform="\1count">;

c =countwords(inputfile);
```

mapper.swift:

```
type filee;

type count;

app (count t) countwords(filee f) {

wc "-w" @filename(f) stdout=@filename(t);

}

string inputNames = "mapper1.txt";

string outputNames = "mapper1.count ";

filee inputfiles[] <mapper;files=inputNames>;

count outputfiles[] <mapper;files=outputNames>;

outputfiles[0] = countwords(inputfiles[0]);
```

wordcount.swift:

```
type filee ;
```

```
type count;
```

```
app (count t) countwords (filee f) {
```

```
  wc "-w" @filename(f) stdout=@filename(t);
```

```
}
```

```
string inputName = "wiki10gb.txt";
```

```
filee inputs[] <mapper;files=inputName>;
```

```
foreach f in inputs {
```

```
  count c<regexexpressionmapper;source=@f, match="(.*?)txt",transform="//1count">;
```

```
  c = countwords(f);
```

```
}
```

MPI Sorting:

```
#include <iostream>
```

```
#include <fstream>
```

```
#include <string>
```

```
#include <algorithm>
```

```
#include <boost/mpi.hpp>
```

```
#include <boost/algorithm/string.hpp>
```

```
#include <boost/serialization/map.hpp>
```

```
using namespace std;
```

```
namespace mpi::boost::mpi;
```

```
void print_help()
```

```
{
```

```
    cout << "Usage: mpiexec word-count <path to file>"<< endl;
```

```
}
```

```
char separator[] = " ,\n-+;:!?()\t[]{}<>`\"";
```

```
bool isseparator (char c) {
```

```
    char* e = separator + sizeof(separator) / sizeof(separator[0]);
```

```
    char* pos = std::find(separator, e, c);
```

```
    return (pos != e);
```

```
}
```

```

struct sort_reverse {
    bool operator()(const std::pair<string,int> &left, const std::pair<string,int> &right) {
        return left.second > right.second;
    }
};

```

```

int main(int argc, char *argv[])
{
    mpi::environment env(argc, argv);
    mpi::communicator world;

```

```

    if (argc != 2)
    {
        if (world.rank() == 0)
        {
            print_help();
        }
        return 1;
    }

```

```

    int chunksize;

    int mapnodescount = world.size() - 1;

    int masterrank = 0;

    int chunk_size_max;

    unsigned int maxoutputlines = 25;

```

```
//masterrank vars

char* buf;

int buf_len;

vector<map<string, unsigned long long int>> stats;

map<string, unsigned long long int> result;


//slaves vars

map<string, unsigned long long int> stat;

// vector<string, unsigned long long int> stat_plain;

char* input = NULL;


if (world.rank() == masterrank)
{
    // read input file

    string word;

    ifstream infile(argv[1]);


    infile.seekg(0, ios::end);

    buf_len = infile.tellg();

    infile.seekg(0, ios::beg);

    buf = new char[buf_len];

    infile.read(buf, buf_len);

    infile.close();
}
```



```

chunksize = buf_len / mapnodescount;

chunk_size_max = chunksize * 2;

mpi::broadcast(world, chunk_size_max, masterrank);

// split in chunks and send
int i, start_index = 0;
for (i = 1; i < mapnodescount; i++){
    if (start_index >= buf_len)
    {
        world.send(i, 0, buf + buf_len, 1);
    }
    else
    {
        int size = chunksize;
        while (size < chunk_size_max && start_index + size <= buf_len &&
            !isseparator(buf[start_index + size - 1]))
        {
            size++;
        }
        buf[start_index + size - 1] = 0;
        world.send(i, 0, buf + start_index, max(size, 1));
        start_index += size;
    }
}

```

```

        world.send(i, 0, buf + min(start_index, buf_len), max(buf_len - start_index + 1, 1));
    }
else {
    // get inputs

    mpi::broadcast(world, chunk_sizemax, masterrank);

    input = new char[chunk_sizemax];

    world.recv(masterrank, 0, input, chunk_sizemax);

    //cout << "worker #" << world.rank() << ": " << input << "" << endl; //DEBUG
}

//barrier and timing
world.barrier();
mpi::timer timer;

// do word count
if(world.rank() != masterrank) {
    char * word;

    word = strtok(input,separator);

    while (word != NULL)
    {
        string s_word(word);

        boost::algorithm::to_lower(s_word);

        stat[s_word] += 1;

        word = strtok (NULL, separator);
    }

    vector<string> smth;

```

```

        world.send(masterrank, 1, smth);
    }

    // collect results

    mpi::gather(world, stat, stats, masterrank);

    if(world.rank() == masterrank)
    {
        for(vector<map<string, unsigned long long int>::iterator it = stats.begin(); it != stats.end();
        ++it) {
            for (map<string, unsigned long long int>::iterator vit = it->begin(); vit != it->end();
            ++vit) {
                result[vit->first] += vit->second;
            }
        }
    }

    // output sorted results

    if (world.rank() == masterrank)
    {
        cout << "Time: " << timer.elapsed() << "s" << endl;

        cout << "Total words: " << result.size() << endl;

        vector<pair<string, unsigned long long int>> output;

        for (map<string, unsigned long long int>::iterator it = result.begin(); it != result.end(); ++it)
        {
            output.push_back(make_pair(it->first, it->second));
        }
    }
}

```

```

    }

    sort(output.begin(), output.end(), sort_reverse());

    for (unsigned int i = 0; i < output.size() && i < maxoutputlines; i++) {
        cout << output[i].first << " => " << output[i].second << endl;
    }
}

return 0;
}

```

Config:

```

#####

## StarCluster Configuration File ##

#####

[global]

DEFAULT_TEMPLATE=akshara

#####

## AWS Credentials and Connection Settings ##

#####

[aws info]

AWS_ACCESS_KEY_ID =AKIAI2QU45XPVS6SEAQQ

AWS_SECRET_ACCESS_KEY =GjbFZIf4xks9oQeWzFfl lkb3g74jxtz2414K6vfj
#your_secret_access_key

AWS_USER_ID= akshara #your userid

```

```
AWS_REGION_NAME = eu-west-2a
```

```
AWS_REGION_HOST=eu-west-2a.amazonaws.com
```

```
## Defining EC2 Keypairs ##
```

```
[key akshu.rsa]
```

```
KEY_LOCATION=/home/ubuntu/akshu.rsa
```

```
## Defining Cluster Templates ##
```

```
$ starcluster start akshara
```

```
[cluster akshara]
```

```
KEYNAME = akshu.rsa
```

```
CLUSTER_SIZE = 2
```

```
CLUSTER_USER = sgeadmin
```

```
CLUSTER_SHELL = bash
```

```
# Uncomment to prepent the cluster tag to the dns name of all nodes created
```

```
# using this cluster config. ie: mycluster-master and mycluster-node001
```

```
# If you choose to enable this option, it's recommended that you enable it in
```

```
# the DEFAULT_TEMPLATE so all nodes will automatically have the prefix
```

```
# DNS_PREFIX = True
```

```
#AMI to use for cluster nodes. These AMIs are for the us-east-1 region.

# Use the 'listpublic' command to list StarCluster AMIs in other regions

# The base i386 StarCluster AMI is ami-9bf9c9f2

# The base x86_64 StarCluster AMI is ami-3393a45a

# The base HVM StarCluster AMI is ami-6b211202

NODE_IMAGE_ID = ami-6b211202

# instance type for all cluster nodes

# (options: m3.large, i2.8xlarge, c3.2xlarge, hs1.8xlarge, c1.xlarge, r3.4xlarge, g2.2xlarge,
m1.small, c1.medium, m3.2xlarge, c3.8xlarge, m2.xlarge, r3.2xlarge, t1.micro, cr1.8xlarge,
r3.8xlarge, cc1.4xlarge, m1.medium, r3.large, c3.xlarge, i2.xlarge, m3.medium, cc2.8xlarge,
m1.large, cg1.4xlarge, i2.2xlarge, c3.large, i2.4xlarge, c3.4xlarge, r3.xlarge, m1.xlarge,
hi1.4xlarge, m2.4xlarge, m2.2xlarge, m3.xlarge)

NODE_INSTANCE_TYPE = c3.large

# Launch cluster in a VPC subnet (OPTIONAL)

#SUBNET_ID=subnet-999999999

# Uncomment to assign public IPs to cluster nodes (VPC-ONLY) (OPTIONAL)

# WARNING: Using public IPs with a VPC requires:

# 1. An internet gateway attached to the VPC

# 2. A route table entry linked to the VPC's internet gateway and associated

# with the VPC subnet with a destination CIDR block of 0.0.0.0/0

# WARNING: Public IPs allow direct access to your VPC nodes from the internet

#PUBLIC_IPS=True

# Uncomment to disable installing/configuring a queueing system on the

# cluster (SGE)

#DISABLE_QUEUE=True

# Uncomment to specify a different instance type for the master node (OPTIONAL)

# (defaults to NODE_INSTANCE_TYPE if not specified)
```

```
MASTER_INSTANCE_TYPE = c3.large

# Uncomment to specify a separate AMI to use for the master node. (OPTIONAL)

# (defaults to NODE_IMAGE_ID if not specified)

MASTER_IMAGE_ID = ami-6b211202

# availability zone to launch the cluster in (OPTIONAL)

# (automatically determined based on volumes (if any) or

# selected by Amazon if not specified)

#AVAILABILITY_ZONE = us-east-1c

# list of volumes to attach to the master node (OPTIONAL)

# these volumes, if any, will be NFS shared to the worker nodes

# see "Configuring EBS Volumes" below on how to define volume sections

#VOLUMES = oceandata, biodata

# list of plugins to load after StarCluster's default setup routines (OPTIONAL)

# see "Configuring StarCluster Plugins" below on how to define plugin sections

#PLUGINS = myplugin, myplugin2

# list of permissions (or firewall rules) to apply to the cluster's security

# group (OPTIONAL).

#PERMISSIONS = ssh, http

# Uncomment to always create a spot cluster when creating a new cluster from

# this template. The following example will place a $0.50 bid for each spot

# request.

#SPOT_BID = 0.50

# Uncomment to specify one or more userdata scripts to use when launching

# cluster instances. Supports cloudinit. All scripts combined must be less than

# 16KB
```

```
#USERDATA_SCRIPTS = /path/to/script1, /path/to/script2
```

```
#####
```

```
## Defining Additional Cluster Templates ##
```

```
#####
```

```
# You can also define multiple cluster templates. You can either supply all  
# configuration options as with smallcluster above, or create an  
# EXTENDS=<cluster_name> variable in the new cluster section to use all  
# settings from <cluster_name> as defaults. Below are example templates that  
# use the EXTENDS feature:
```

```
# [cluster mediumcluster]
```

```
# Declares that this cluster uses smallcluster as defaults
```

```
# EXTENDS=smallcluster
```

```
# This section is the same as smallcluster except for the following settings:
```

```
# KEYNAME=myotherkey
```

```
# NODE_INSTANCE_TYPE = c1.xlarge
```

```
# CLUSTER_SIZE=8
```

```
# VOLUMES = biodata2
```

```
# [cluster largecluster]
```

```
# Declares that this cluster uses mediumcluster as defaults
```

```
# EXTENDS=mediumcluster
```

```
# This section is the same as mediumcluster except for the following variables:
```

```
# CLUSTER_SIZE=16
```



```
#####
```

```
## Configuring EBS Volumes ##
```

```
#####
```

```
# StarCluster can attach one or more EBS volumes to the master and then
# NFS_share these volumes to all of the worker nodes. A new [volume] section
# must be created for each EBS volume you wish to use with StarCluster. The
# section name is a tag for your volume. This tag is used in the VOLUMES
# setting of a cluster template to declare that an EBS volume is to be mounted
# and nfs shared on the cluster. (see the commented VOLUMES setting in the
# example 'smallcluster' template above) Below are some examples of defining
# and configuring EBS volumes to be used with StarCluster:
```

```
# Sections starting with "volume" define your EBS volumes
```

```
# [volume biodata]
```

```
# attach vol-c9999999 to /home on master node and NFS-shre to worker nodes
```

```
# VOLUME_ID = vol-c9999999
```

```
# MOUNT_PATH = /home
```

```
# Same volume as above, but mounts to different location
```

```
# [volume biodata2]
```

```
# VOLUME_ID = vol-c9999999
```

```
# MOUNT_PATH = /opt/
```

```
# Another volume example
```

```
# [volume oceandata]

# VOLUME_ID = vol-d7777777

# MOUNT_PATH = /mydata


# By default StarCluster will attempt first to mount the entire volume device,
# failing that it will try the first partition. If you have more than one
# partition you will need to set the PARTITION number, e.g.:

# [volume oceandata]

# VOLUME_ID = vol-d7777777

# MOUNT_PATH = /mydata

# PARTITION = 2


#####

## Configuring Security Group Permissions ##

#####

# Sections starting with "permission" define security group rules to
# automatically apply to newly created clusters. IP_PROTOCOL in the following
# examples can be: tcp, udp, or icmp. CIDR_IP defaults to 0.0.0.0/0 or
# "open to the world"


# open port 80 on the cluster to the world

# [permission http]

# IP_PROTOCOL = tcp

# FROM_PORT = 80

# TO_PORT = 80
```

```
# open https on the cluster to the world
```

```
# [permission https]
```

```
# IP_PROTOCOL = tcp
```

```
# FROM_PORT = 443
```

```
# TO_PORT = 443
```

```
# open port 80 on the cluster to an ip range using CIDR_IP
```

```
# [permission http]
```

```
# IP_PROTOCOL = tcp
```

```
# FROM_PORT = 80
```

```
# TO_PORT = 80
```

```
# CIDR_IP = 18.0.0.0/8
```

```
# restrict ssh access to a single ip address (<your_ip>)
```

```
# [permission ssh]
```

```
# IP_PROTOCOL = tcp
```

```
# FROM_PORT = 22
```

```
# TO_PORT = 22
```

```
# CIDR_IP = <your_ip>/32
```

```
#####
```

```
## Configuring StarCluster Plugins ##
```

```
#####
```

```
# Sections starting with "plugin" define a custom python class which perform
# additional configurations to StarCluster's default routines. These plugins
# can be assigned to a cluster template to customize the setup procedure when
# starting a cluster from this template (see the commented PLUGINS setting in
# the 'smallcluster' template above). Below is an example of defining a user
# plugin called 'myplugin':
```

```
# [plugin myplugin]
```

```
# NOTE: myplugin module must either live in ~/.starcluster/plugins or be
```

```
# on your PYTHONPATH
```

```
# SETUP_CLASS = myplugin.SetupClass
```

```
# extra settings are passed as __init__ arguments to your plugin:
```

```
# SOME_PARAM_FOR_MY_PLUGIN = 1
```

```
# SOME_OTHER_PARAM = 2
```

```
#####
```

```
## Built-in Plugins ##
```

```
#####
```

```
# The following plugins ship with StarCluster and should work out-of-the-box.
```

```
# Uncomment as needed. Don't forget to update your PLUGINS list!
```

```
# See http://star.mit.edu/cluster/docs/latest/plugins for plugin details.
```

```
#
```

```
# Use this plugin to install one or more packages on all nodes
```

```
# [plugin pkginstaller]
```

```
# SETUP_CLASS = starcluster.plugins.pkginstaller.PackageInstaller
```

```
# # list of apt-get installable packages

# PACKAGES = mongodb, python-pymongo

#

# Use this plugin to create one or more cluster users and download all user ssh

# keys to $HOME/.starcluster/user_keys/<cluster>-<region>.tar.gz

# [plugin createusers]

# SETUP_CLASS = starcluster.plugins.users.CreateUsers

# NUM_USERS = 30

# # you can also comment out NUM_USERS and specify exact usernames, e.g.

# # usernames = linus, tux, larry

# DOWNLOAD_KEYS = True

#

# Use this plugin to configure the Condor queueing system

# [plugin condor]

# SETUP_CLASS = starcluster.plugins.condor.CondorPlugin

#

# The SGE plugin is enabled by default and not strictly required. Only use this

# if you want to tweak advanced settings in which case you should also set

# DISABLE_QUEUE=TRUE in your cluster template. See the plugin doc for more

# details.

# [plugin sge]

# SETUP_CLASS = starcluster.plugins.sge.SGEPlugin

# MASTER_IS_EXEC_HOST = False

#

# The IPcluster plugin configures a parallel IPython cluster with optional
```

```
# web notebook support. This allows you to run Python code in parallel with low
# latency message passing via ZeroMQ.

# [plugin ipcluster]

# SETUP_CLASS = starcluster.plugins.ipcluster.IPCluster

# # Enable the IPython notebook server (optional)

# ENABLE_NOTEBOOK = True

# # Set a password for the notebook for increased security

# # This is optional but *highly* recommended

# NOTEBOOK_PASSWD = a-secret-password

# # Set a custom directory for storing/loading notebooks (optional)

# NOTEBOOK_DIRECTORY = /path/to/notebook/dir

# # Set a custom packer. Must be one of 'json', 'pickle', or 'msgpack'

# # This is optional.

# PACKER = pickle

#

# Use this plugin to create a cluster SSH "dashboard" using tmux. The plugin
# creates a tmux session on the master node that automatically connects to all
# the worker nodes over SSH. Attaching to the session shows a separate window
# for each node and each window is logged into the node via SSH.

# [plugin tmux]

# SETUP_CLASS = starcluster.plugins.tmux.TmuxControlCenter

#

# Use this plugin to change the default MPI implementation on the
# cluster from OpenMPI to MPICH2.

# [plugin mpich2]
```

```
# SETUP_CLASS = starcluster.plugins.mpich2.MPICH2Setup
#
# Configure a hadoop cluster. (includes dumbo setup)
# [plugin hadoop]
# SETUP_CLASS = starcluster.plugins.hadoop.Hadoop
#
# Configure a distributed MySQL Cluster
# [plugin mysqlcluster]
# SETUP_CLASS = starcluster.plugins.mysql.MysqlCluster
# NUM_REPLICAS = 2
# DATA_MEMORY = 80M
# INDEX_MEMORY = 18M
# DUMP_FILE = test.sql
# DUMP_INTERVAL = 60
# DEDICATED_QUERY = True
# NUM_DATA_NODES = 2
#
# Install and setup an Xvfb server on each cluster node
# [plugin xvfb]
# SETUP_CLASS = starcluster.plugins.xvfb.XvfbSetup
```